



# Data-driven and Learning-based Control

## Policy optimization

Erica Salvato



**UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE**



# Table of Contents

1 A brief recap

► A brief recap

► Policy Optimization

► Actor Critic

► Conclusion



# What we know so far?

## 1 A brief recap

We introduced **Reinforcement Learning**:

- as a discrete-time, stochastic or deterministic optimal control problem
- where system's dynamical model is unknown → **Model-free control**
- and the only essential requirements are the **reward function** and data acquired directly from the real system.

$$\begin{aligned}\pi^* &= \arg \max_{\pi} \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k h \left( \mathbf{x}^{(k)}, \pi \left( \mathbf{x}^{(k)} \right), \mathbf{x}^{(k+1)} \right) \right] \\ &= \arg \max_{\pi} \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r^{(k+1)} \right]\end{aligned}$$



# What we know so far?

## 1 A brief recap

We observed that we can classify Reinforcement Learning approaches in:

- **Value-function methods:**

The policy is implicitly defined via  $V(x^{(k)})$  or  $Q(x^{(k)}, u^{(k)})$  directly relying on Bellman's equations



# What we know so far?

## 1 A brief recap

We observed that we can classify Reinforcement Learning approaches in:

- **Value-function methods:**

The policy is implicitly defined via  $V(x^{(k)})$  or  $Q(x^{(k)}, u^{(k)})$  directly relying on Bellman's equations

- **Policy optimization methods:**

The policy is a parameterized function whose weights are learned to maximize the expected cumulative discounted reward



# What we know so far?

## 1 A brief recap

We observed that we can classify Reinforcement Learning approaches in:

- **Value-function methods:**

The policy is implicitly defined via  $V(x^{(k)})$  or  $Q(x^{(k)}, u^{(k)})$  directly relying on Bellman's equations → **Critic**

- **Policy optimization methods:**

The policy is a parameterized function whose weights are learned to maximize the expected cumulative discounted reward → **Actor**



# What we know so far?

## 1 A brief recap

We observed that we can classify Reinforcement Learning approaches in:

- **Value-function methods:**

The policy is implicitly defined via  $V(x^{(k)})$  or  $Q(x^{(k)}, u^{(k)})$  directly relying on Bellman's equations → **Critic**

- **Policy optimization methods:**

The policy is a parameterized function whose weights are learned to maximize the expected cumulative discounted reward → **Actor**

- **Actor-critic methods:**

Merging the two ideas by guiding the actor's learning based on the critic's estimated return



# What we know so far?

## 1 A brief recap

We studied and applied **Value-function methods**.

We observed that they can be classified:

- Depending on the evaluation procedure in:

- On-policy algorithms

- SARSA

- Off-policy algorithms

- Q-learning

- Depending on the state space representation

- Tabular

- SARSA
- Q-learning

- Linear function approximation

- SARSA
- Q-learning

- Non-linear function approximation

- Deep Q-learning

Here, we assumed to work with a **discrete compact action set**  $\mathcal{U}$ .





# What we know so far?

## 1 A brief recap

We studied and applied **Value-function methods**.

We observed that they can be classified:

- Depending on the evaluation procedure in:

- On-policy algorithms

- SARSA

- Off-policy algorithms

- Q-learning

- Depending on the state space representation

- Tabular

- SARSA
- Q-learning

- Linear function approximation

- SARSA
- Q-learning

- Non-linear function approximation

- Deep Q-learning

Here, we assumed to work with a **discrete compact action set**  $\mathcal{U}$ . → **What if we can't?**



# Table of Contents

## 2 Policy Optimization

- ▶ A brief recap
- ▶ **Policy Optimization**
- ▶ Actor Critic
- ▶ Conclusion



# Policy-Based methods

## 2 Policy Optimization

Policy-based methods aim to learn directly parametrized policy

$$\pi_{\theta}(u|x) = \Pr(u^{(k)}|x^{(k)}, \theta)$$

- Advantages:
  - Allow to work also with high-dimensional or continuous action spaces
  - Learn stochastic policies
- Disadvantages:
  - Typically converge to a local rather than global optimum
  - Evaluating a policy is typically inefficient and high variance



# Epsilon-Greedy vs. Stochastic Policies

## 2 Policy Optimization

- **Epsilon-Greedy Policy**

- Exploit (choose best action) with probability  $1 - \epsilon$
- Explore (choose random action) with probability  $\epsilon$
- Stochastic component: when to explore and exploration
- Deterministic component: exploitation

**N.B.** Even when it explores it chooses equally among all actions, i.e., the probability of choosing the worst-appearing action is equal to that of choosing the next-to-best action.



# Epsilon-Greedy vs. Stochastic Policies

## 2 Policy Optimization

- **Epsilon-Greedy Policy**

- Exploit (choose best action) with probability  $1 - \epsilon$
- Explore (choose random action) with probability  $\epsilon$
- Stochastic component: when to explore and exploration
- Deterministic component: exploitation

**N.B.** Even when it explores it chooses equally among all actions, i.e., the probability of choosing the worst-appearing action is equal to that of choosing the next-to-best action.

- **Stochastic Policy**

- Probabilistic action selection
- Probability distribution over actions



# Policy-Based methods

## 2 Policy Optimization

- Assume working with an episodic approach
- In an episode the RL controller interacts with the system and collects a trajectory

$$\tau = \left( x^{(0)}, u^{(0)}, x^{(1)}, u^{(1)}, \dots, x^{(H)}, u^{(H)} \right)$$

- We define the trajectory distribution

$$p(\tau) = p(x^{(0)}) \prod_{k=0}^{H-1} \pi_{\theta}(u^{(k)} | x^{(k)}) T(x^{(k+1)} | x^{(k)}, u^{(k)})$$

- The RL objective can be then expressed as an expectation under the trajectory distribution

$$\pi_{\theta}^* = \arg \max_{\pi_{\theta}} \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} \gamma^i r^{(i+1)} \right]$$



# Policy-Based methods

## 2 Policy Optimization

The goal is to find the optimal parameter vector  $\theta^* \in \mathbb{R}^t$  such that

$$\theta^* = \arg \max_{\pi_{\theta}} \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} \gamma^i r^{(i+1)} \right]$$

To simplify the notation assume  $\gamma^i r^{(i+1)} = R(x^{(i)}, u^{(i)})$ , therefore:

$$\theta^* = \arg \max_{\pi_{\theta}} \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right]$$



# Policy-Based methods

## 2 Policy Optimization

The goal is to find the optimal parameter vector  $\theta^* \in \mathbb{R}^t$  such that

$$\theta^* = \arg \max_{\pi_{\theta}} \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} \gamma^i r^{(i+1)} \right]$$

To simplify the notation assume  $\gamma^i r^{(i+1)} = R(x^{(i)}, u^{(i)})$ , therefore:

$$\theta^* = \arg \max_{\pi_{\theta}} \mathbb{E}_{\tau \sim p(\tau)} \underbrace{\left[ \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right]}_{J(\theta)}$$





# Policy-Based methods

## 2 Policy Optimization

Therefore we have to optimize  $J(\theta)$  with respect to  $\theta$

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right]$$



# Policy-Based methods

## 2 Policy Optimization

Therefore we have to optimize  $J(\theta)$  with respect to  $\theta$

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right]$$

The procedure is as follows:

1. Estimate the gradient  $\nabla_{\theta} J(\theta)$



# Policy-Based methods

## 2 Policy Optimization

Therefore we have to optimize  $J(\theta)$  with respect to  $\theta$


$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right]$$

The procedure is as follows:

1. Estimate the gradient  $\nabla_{\theta} J(\theta)$
2. Cast the learning process as approximate gradient ascent on  $J(\theta)$

$$\theta = \theta + \alpha \nabla_{\theta} J(\theta) = \theta + \alpha \begin{bmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_t} \end{bmatrix}$$

where  $\alpha$  is a step-size parameter





# Direct policy gradient

## 2 Policy Optimization

How can we compute  $\nabla_{\theta} J(\theta)$ ?

- to ease the notation define  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$ . therefore

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} [r(\tau)] = \int p_{\theta}(\tau) r(\tau) d\tau$$

- the gradient therefore can be written as follows:

$$\nabla_{\theta} J(\theta) = \int \nabla_{\theta} p_{\theta}(\tau) r(\tau) d\tau$$



# Direct policy gradient

## 2 Policy Optimization

How can we compute  $\nabla_{\theta} J(\theta)$ ?

- to ease the notation define  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$ . therefore

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} [r(\tau)] = \int p_{\theta}(\tau) r(\tau) d\tau$$

- the gradient therefore can be written as follows:

$$\nabla_{\theta} J(\theta) = \int p_{\theta}(\tau) \frac{\nabla_{\theta} p_{\theta}(\tau)}{p_{\theta}(\tau)} r(\tau) d\tau$$



# Direct policy gradient

## 2 Policy Optimization

How can we compute  $\nabla_{\theta} J(\theta)$ ?

- to ease the notation define  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$ . therefore

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} [r(\tau)] = \int p_{\theta}(\tau) r(\tau) d\tau$$

- the gradient therefore can be written as follows:

$$\nabla_{\theta} J(\theta) = \int p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) r(\tau) d\tau$$



# Direct policy gradient

## 2 Policy Optimization

How can we compute  $\nabla_{\theta} J(\theta)$ ?

- to ease the notation define  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$ . therefore

$$J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} [r(\tau)] = \int p_{\theta}(\tau) r(\tau) d\tau$$

- the gradient therefore can be written as follows:

$$\begin{aligned} \nabla_{\theta} J(\theta) &= \int p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) r(\tau) d\tau \\ &= \mathbb{E}_{\tau \sim p(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) r(\tau)] \end{aligned}$$



# Direct policy gradient

## 2 Policy Optimization

Recall that

$$p(\tau) = p(x^{(0)}) \prod_{k=0}^{H-1} \pi_{\theta}(u^{(k)} | x^{(k)}) T(x^{(k+1)} | x^{(k)}, u^{(k)})$$

Then

$$\log p(\tau) = \log p(x^{(0)}) + \sum_{k=0}^{H-1} \log \pi_{\theta}(u^{(k)} | x^{(k)}) + \log T(x^{(k+1)} | x^{(k)}, u^{(k)})$$

and therefore:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) r(\tau)] = \mathbb{E}_{\tau \sim p(\tau)} \left[ \nabla_{\theta} \left[ \sum_{k=0}^{H-1} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \right] r(\tau) \right]$$





# Direct policy gradient

## 2 Policy Optimization

Finally by substituting  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$  we obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

where everything inside the expectation is known.



# Direct policy gradient

## 2 Policy Optimization

Finally by substituting  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$  we obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

where everything inside the expectation is known. Now, supposing to sample  $N$  trajectories we can write the expectation as follows:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u_j^{(k)} | x_j^{(k)}) \left( \sum_{i=0}^{H-1} R(x_j^{(i)}, u_j^{(i)}) \right) \right]$$



# Direct policy gradient

## 2 Policy Optimization

Finally by substituting  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$  we obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

where everything inside the expectation is known. Now, supposing to sample  $N$  trajectories we can write the expectation as follows:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{j=1}^N \left[ \underbrace{\sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u_j^{(k)} | x_j^{(k)})}_{\text{maximum likelihood estimation}} \left( \sum_{i=0}^{H-1} R(x_j^{(i)}, u_j^{(i)}) \right) \right]$$



# Direct policy gradient

## 2 Policy Optimization

Finally by substituting  $r(\tau) = \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)})$  we obtain

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

where everything inside the expectation is known. Now, supposing to sample  $N$  trajectories we can write the expectation as follows:

$$\nabla_{\theta} J(\theta) \approx \frac{1}{N} \sum_{j=1}^N \left[ \underbrace{\sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u_j^{(k)} | x_j^{(k)})}_{\text{maximum likelihood estimation}} \underbrace{\left( \sum_{i=0}^{H-1} R(x_j^{(i)}, u_j^{(i)}) \right)}_{\text{weighted maximum likelihood estimation}} \right]$$



# Softmax Policy

## 2 Policy Optimization

### Softmax Policy

The softmax policy is a stochastic policy that selects a control input  $u^{(k)}$  according to:

$$\Pr(u^{(k)} | x^{(k)}, \theta) = \frac{e^{\phi(x^{(k)}, u^{(k)})^\top \theta}}{\sum_i e^{\phi(x^{(k)}, u_i)^\top \theta}}$$

where:

- $\phi(x^{(k)}, u^{(k)})^\top \theta$  is the approximated action-value function.

It is used, again, when actions belong to a discrete and compact set.



# Gaussian Policy

## 2 Policy Optimization

### Gaussian Policy

The Gaussian policy is a stochastic policy that models the probability distribution over actions using a Gaussian (normal) distribution:

$$\Pr(u^{(k)} | x^{(k)}, \theta) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(u^{(k)} - \mu)^2}{2\sigma^2}\right)$$

where:

- $\mu = \phi(x^{(k)})^\top \theta$  is the approximated value function.
- $\sigma$  is the standard deviation

It is used, instead, in the case of continuous action space.



# REINFORCE algorithm

## 2 Policy Optimization

1. Initialize  $\pi_{\theta_l} = \pi_{\theta_0}$
2. Sample  $N$  trajectories  $\tau_i, i = 1, \dots, N$  by running  $\pi_{\theta_l}$  on the environment
3. Evaluate the policy gradient

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=0}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

4. Take a gradient step update

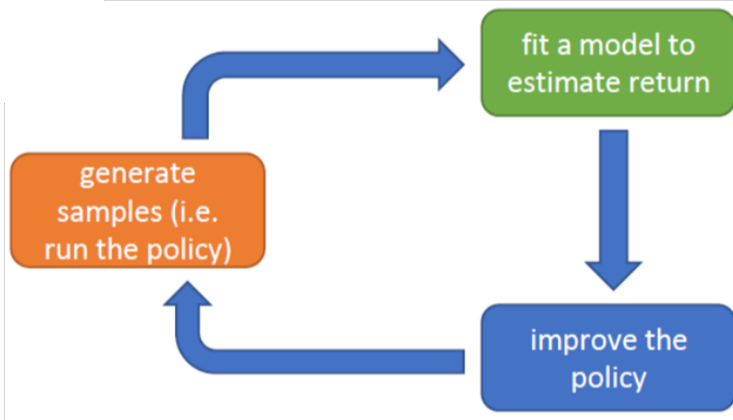
$$\theta_{l+1} = \theta_l + \alpha \nabla_{\theta_l} J(\theta_l)$$

5. Repeat from step 2.



# REINFORCE algorithm

## 2 Policy Optimization







# Challenges of Policy Gradient Methods

## 2 Policy Optimization

Consider the policy gradient evaluation formula:

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

- It requires sampling entire trajectories before each gradient update



# Challenges of Policy Gradient Methods

## 2 Policy Optimization

Consider the policy gradient evaluation formula:

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

- It requires sampling entire trajectories before each gradient update → **extensive sampling**



# Challenges of Policy Gradient Methods

## 2 Policy Optimization

Consider the policy gradient evaluation formula:

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

- It requires sampling entire trajectories before each gradient update → **extensive sampling**
- Sampling multiple trajectories from an untrained policy leads to highly variable behaviors



# Challenges of Policy Gradient Methods

## 2 Policy Optimization

Consider the policy gradient evaluation formula:


$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

- It requires sampling entire trajectories before each gradient update → **extensive sampling**
- Sampling multiple trajectories from an untrained policy leads to highly variable behaviors → **high variance**



# High Variance Policy Gradient Methods

## 2 Policy Optimization




- Depending on the sample, the policy gradient can vary wildly
- This negatively affects learning: worse performance, slower convergence



# High Variance Policy Gradient Methods

## 2 Policy Optimization




- Depending on the sample, the policy gradient can vary wildly
- This negatively affects learning: worse performance, slower convergence



# High Variance Policy Gradient Methods

## 2 Policy Optimization




- Depending on the sample, the policy gradient can vary wildly
- This negatively affects learning: worse performance, slower convergence



# High Variance Policy Gradient Methods

## 2 Policy Optimization




- Depending on the sample, the policy gradient can vary wildly
- This negatively affects learning: worse performance, slower convergence





# High Variance Policy Gradient Methods

## 2 Policy Optimization



- Depending on the sample, the policy gradient can vary wildly
- This negatively affects learning: worse performance, slower convergence



# Reduce Variance of Policy Gradient Methods

## 2 Policy Optimization

1. A first simple approach to reduce the variance entails using causality:

*policy at time  $l$  cannot affect reward at time  $i < l$*

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=k}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$



# Reduce Variance of Policy Gradient Methods

## 2 Policy Optimization

2. A second (more important) approach to reduce the variance introduces the concept of **baseline**

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=1}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( x_j^{(k)}, u_j^{(k)} \right) \left( \sum_{i=1}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) - b \right) \right]$$

Intuitively it allows to “center” our returns, such that:

- behavior better than average gets increased
- behavior worse than average gets decreased



# Reduce Variance of Policy Gradient Methods

## 2 Policy Optimization

Notice that adding the baseline does not change the value of the expected gradient

$$\begin{aligned}\mathbb{E}_{\tau \sim p(\tau)} [\nabla_{\theta} \log p_{\theta}(\tau) b] &= \int p_{\theta}(\tau) \nabla_{\theta} \log p_{\theta}(\tau) b d\tau = \int \nabla_{\theta} p_{\theta}(\tau) b d\tau \\ &= b \nabla_{\theta} \int p_{\theta}(\tau) d\tau = b \nabla_{\theta} 1 = 0\end{aligned}$$

thus making our estimate of the gradient (with baseline) unbiased in expectation.


An extremely effective choice of baseline is the average return:

$$\frac{1}{N} \sum_{i=1}^N r(\tau)$$



# Policy Gradient with and without baseline


## 2 Policy Optimization





# Policy Gradient with and without baseline

## 2 Policy Optimization



Is the policy gradient on-policy or off-policy?




# Sample inefficiency of Policy Gradient Methods

## 2 Policy Optimization

Requiring a complete trajectory to compute

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

indicates our reliance on a Monte Carlo procedure once more.






# Sample inefficiency of Policy Gradient Methods

## 2 Policy Optimization

Requiring a complete trajectory to compute

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) \left( \sum_{i=0}^{H-1} R(x^{(i)}, u^{(i)}) \right) \right]$$

indicates our reliance on a Monte Carlo procedure once more.



The following main drawbacks arise:

- It requires a large number of random samples to achieve accurate results → **high computational costs**
- As the dimensionality of the problem increases, the number of samples required for accurate estimates grows exponentially → **curse of dimensionality**





# Reduce Extensive Sampling of Policy Gradient Methods

## 2 Policy Optimization

A solution might be to substitute  $\sum_{i=k}^{H-1} R(x^{(i)}, u^{(i)})$  with the estimate action value function  $Q(x^{(k)}, u^{(k)})$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) Q(x^{(k)}, u^{(k)}) \right]$$

Hence, by applying what we already studied we can use a **Critic** to estimate the action-value function:

$$Q_w(x^{(k)}, u^{(k)}) \approx Q(x^{(k)}, u^{(k)})$$

In this way, the Critic update can be performed in a TD learning fashion.



# Reduce Extensive Sampling of Policy Gradient Methods

## 2 Policy Optimization

A solution might be to substitute  $\sum_{i=k}^{H-1} R(x^{(i)}, u^{(i)})$  with the estimate action value function  $Q(x^{(k)}, u^{(k)})$

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta}(u^{(k)} | x^{(k)}) Q(x^{(k)}, u^{(k)}) \right]$$

Hence, by applying what we already studied we can use a **Critic** to estimate the action-value function:

$$Q_w(x^{(k)}, u^{(k)}) \approx Q(x^{(k)}, u^{(k)})$$

In this way, the Critic update can be performed in a TD learning fashion.



**Actor Critic**



# Table of Contents

## 3 Actor Critic

- ▶ A brief recap
- ▶ Policy Optimization
- ▶ Actor Critic
- ▶ Conclusion



# Actor Critic Methods

## 3 Actor Critic

Actor-critic algorithms maintain two sets of parameters:

- **Critic** updates action-value function parameters  $w$
- **Actor** updates policy parameters  $\theta$  in the direction suggested by critic

thus following an approximate policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta} \left( u^{(k)} | x^{(k)} \right) Q_w \left( x^{(k)}, u^{(k)} \right) \right]$$



# Critic Role


## 3 Actor Critic

The Critic is performing nothing more than policy evaluation:

*Given  $\theta$  how good is policy  $\pi_\theta$ ?*

We already know that this operation can be performed:

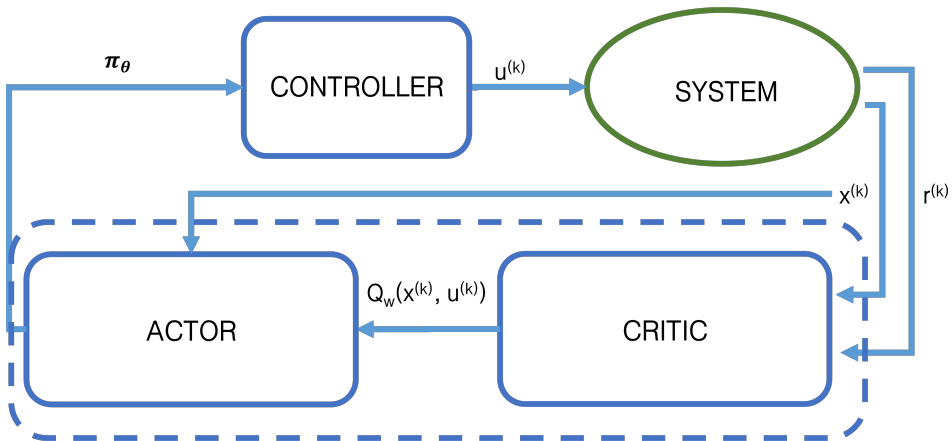
- via Monte-Carlo policy evaluation
- via Temporal-Difference learning





# Actor Critic Scheme

## 3 Actor Critic





# Actor Critic Algorithm with SARSA critic

## 3 Actor Critic

**Initialization.**  $\theta, w$  choose  $\phi(x, u)$  and  $\pi_\theta$

**Repeat** (for each episode)

- Set  $x^{(0)}$
- Select an action  $u^{(k)}$  with  $\pi_\theta$
- Repeat for each step of the episode until the terminal condition is met
  - Perform the action  $u^{(k)}$ , observe  $x^{(k+1)}$  and  $r^{(k+1)}$
  - Select an action  $u^{(k+1)}$  with  $\pi_\theta$
  - $Q(x^{(k)}, u^{(k)}) = w^\top \phi(x^{(k)}, u^{(k)})$  and  $Q(x^{(k+1)}, u^{(k+1)}) = w^\top \phi(x^{(k+1)}, u^{(k+1)})$
  - $\theta = \theta + \beta \nabla_\theta \log \pi_\theta(u^{(k)} | x^{(k)}) Q_w(x^{(k)}, u^{(k)})$
  - $w = w + \alpha [r^{(k+1)} + \gamma Q(x^{(k+1)}, u^{(k+1)}) - Q(x^{(k)}, u^{(k)})] \nabla_w Q(x^{(k)}, u^{(k)})$
  - Update  $x^{(k)} = x^{(k+1)}, u^{(k)} = u^{(k+1)}$



# Approximated Policy Gradient

## 3 Actor Critic

The use of approximating functions to calculate the Critic leads to the achievement of approximations of even  $\nabla_{\theta} J(\theta)$ .

- What if the chosen features are not convenient for goal achievement?
- What if the chosen features are completely wrong?

Therefore, we need to carefully choose value function approximation.

There exists a theorem that allows us to provide some guarantees.





# Approximated Policy Gradient

3 Actor Critic

## Compatible Function Approximation Theorem

If the following two conditions are satisfied:

1. Value function approximator is compatible to the policy

$$\nabla_w Q_w(x, u) = \nabla_{\theta} \log \pi_{\theta}(u|x)$$

2. Value function parameters  $w$  minimise the mean-squared error

$$\mathbb{E}_{\pi_{\theta}} \left[ (Q_{\pi_{\theta}}(x, u) - Q_w(x, u))^2 \right]$$

Then the policy gradient is exact:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(u|x) Q_w(x, u)]$$



# Proof of Compatible Function Approximation Theorem

## 3 Actor Critic

If  $w$  is such that the mean-squared error is minimized, then its gradient with respect to  $w$  must be 0:

$$\nabla_w \mathbb{E}_{\pi_\theta} \left[ (Q_{\pi_\theta}(x, u) - Q_w(x, u))^2 \right] = 0$$

$$\mathbb{E}_{\pi_\theta} [(Q_{\pi_\theta}(x, u) - Q_w(x, u)) \nabla_w Q_w(x, u)] = 0$$

Imposing  $\nabla_w Q_w(x, u) = \nabla_\theta \log \pi_\theta(u|x)$

$$\mathbb{E}_{\pi_\theta} [(Q_{\pi_\theta}(x, u) - Q_w(x, u)) \nabla_\theta \log \pi_\theta(u|x)] = 0$$

$$\mathbb{E}_{\pi_\theta} [Q_{\pi_\theta}(x, u) \nabla_\theta \log \pi_\theta(u|x)] = \mathbb{E}_{\pi_\theta} [Q_w(x, u) \nabla_\theta \log \pi_\theta(u|x)]$$

Therefore can be substituted in the policy gradient thus obtaining

$$\nabla_\theta J(\theta) = \mathbb{E}_{\pi_\theta} [\nabla_\theta \log \pi_\theta(u|x) Q_w(x, u)]$$





# Reduce Variance of Actor Critic Methods

## 3 Actor Critic

Also in this case, in order to reduce the variance we can introduce a baseline

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(u|x) (Q_w(x, u) - b)]$$

A good baseline, in this case, is the value function  $b = V_{\pi_{\theta}}(x)$ .



# Reduce Variance of Actor Critic Methods

## 3 Actor Critic

Also in this case, in order to reduce the variance we can introduce a baseline

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(u|x) (Q_w(x, u) - b)]$$

A good baseline, in this case, is the value function  $b = V_{\pi_{\theta}}(x)$ .

Then by defining the advantage function

$$A_{\pi_{\theta}}(x, u) = Q_{\pi_{\theta}}(x, u) - V_{\pi_{\theta}}(x)$$

we can rewrite the policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(u|x) A_{\pi_{\theta}}(x, u)]$$



# Estimating the Advantage Function

## 3 Actor Critic

Therefore, we need to estimate the advantage function.

1. A way can consist of estimating both  $V_{\pi_{\theta}}(x)$  and  $Q_{\pi_{\theta}}(x, u)$  using two different function approximators with two parameter vectors  $v$  and  $w$ .

$$V_v(x) \approx V_{\pi_{\theta}}(x)$$

$$Q_w(x, u) \approx Q_{\pi_{\theta}}(x, u)$$

$$A(x, u) = Q_w(x, u) - V_v(x)$$

And updating both value functions by TD learning for example



# Estimating the Advantage Function

## 3 Actor Critic

2. A second approach can consist in observing that

$$Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) = r^{(k+1)} + \mathbb{E}_{\pi_{\theta}} [V(x^{(k+1)})]$$

Therefore

$$A(x^{(k)}, u^{(k)}) = Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) - V_{\pi_{\theta}}(x^{(k)}) = r^{(k+1)} + V_{\pi_{\theta}}(x^{(k+1)}) - V_{\pi_{\theta}}(x^{(k)})$$



# Estimating the Advantage Function

## 3 Actor Critic

2. A second approach can consist in observing that

$$Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) = r^{(k+1)} + \mathbb{E}_{\pi_{\theta}} [V(x^{(k+1)})]$$

Therefore

$$A(x^{(k)}, u^{(k)}) = Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) - V_{\pi_{\theta}}(x^{(k)}) = \underbrace{r^{(k+1)} + V_{\pi_{\theta}}(x^{(k+1)})}_{\delta_{\pi_{\theta}}^{(k+1)}} - V_{\pi_{\theta}}(x^{(k)})$$



# Estimating the Advantage Function

## 3 Actor Critic

2. A second approach can consist in observing that

$$Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) = r^{(k+1)} + \mathbb{E}_{\pi_{\theta}} [V(x^{(k+1)})]$$

Therefore

$$A(x^{(k)}, u^{(k)}) = Q_{\pi_{\theta}}(x^{(k)}, u^{(k)}) - V_{\pi_{\theta}}(x^{(k)}) = \underbrace{r^{(k+1)} + V_{\pi_{\theta}}(x^{(k+1)}) - V_{\pi_{\theta}}(x^{(k)})}_{\delta_{\pi_{\theta}}^{(k+1)}}$$

Therefore we can use the TD error to compute the policy gradient

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(u|x) \delta_{\pi_{\theta}}^{(k+1)}]$$

thus requiring only one set of parameters (those used to approximate  $V_{\pi_{\theta}}$ )





# Summary

## 3 Actor Critic

The policy gradient has many equivalent forms

$$\nabla_{\theta_l} J(\theta_l) \approx \frac{1}{N} \sum_{j=1}^N \left[ \sum_{k=0}^{H-1} \nabla_{\theta_l} \log \pi_{\theta_l} \left( u_j^{(k)} | x_j^{(k)} \right) \left( \sum_{i=0}^{H-1} R \left( x_j^{(i)}, u_j^{(i)} \right) \right) \right]$$

REINFORCE

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta} \left( u^{(k)} | x^{(k)} \right) Q_w \left( x^{(k)}, u^{(k)} \right) \right]$$

Actor Critic with SARSA critic

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta} \left( u^{(k)} | x^{(k)} \right) A_w \left( x^{(k)}, u^{(k)} \right) \right]$$

Advantage Actor Critic

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim p(\tau)} \left[ \sum_{k=0}^{H-1} \nabla_{\theta} \log \pi_{\theta} \left( u^{(k)} | x^{(k)} \right) \delta \right]$$

TD Actor Critic

- Each one leading to a stochastic gradient ascent algorithm
- Critic uses policy evaluation (e.g. MC or TD learning) to estimate  $Q$ ,  $A$  or  $V$



# Table of Contents

## 4 Conclusion

▶ A brief recap

▶ Policy Optimization

▶ Actor Critic

▶ Conclusion



# Course summary

## 4 Conclusion

- a) We started introducing the **optimal control problems** in the general form, highlighting the roles of **controllers** and **dynamical systems**.
- b) We observed how to **solve optimal problems** in cases where the dynamical system **model is known**
  - Markov Decision Process (MDP)
  - Linear Quadratic Regulator (LQR)
- c) We observed how optimal control problems involving **non-linear dynamics** can be solved **with LQR**
  - Iterative Linear Quadratic Regulator (iLQR)



# Course summary

## 4 Conclusion

- d) We then moved to the so-called **model-free approaches**, where a model of the dynamical system is not needed
  - Monte Carlo (MC)
  - Temporal Difference learning (TD)
  - Reinforcement Learning (RL)
  
- e) We studied how to apply **RL in different settings**:
  - Bellman's equation-based (or **Value function**) approaches: considering different design choices of the state space
    - Tabular
    - Linear Function approximation
    - Non-linear Function approximation (Neural networks)



## Course summary

### 4 Conclusion

- Approaches that learn the policy directly: allowing for different choices of the action space (**Policy Gradient**).
- Hybrid approaches: combining both ideas, thus allowing to work with all possible choices of both state and action spaces (**Actor Critic**).



Questions' time!



**UNIVERSITÀ  
DEGLI STUDI  
DI TRIESTE**